

Lecture IN-2147 Parallel Programming

SoSe 2018

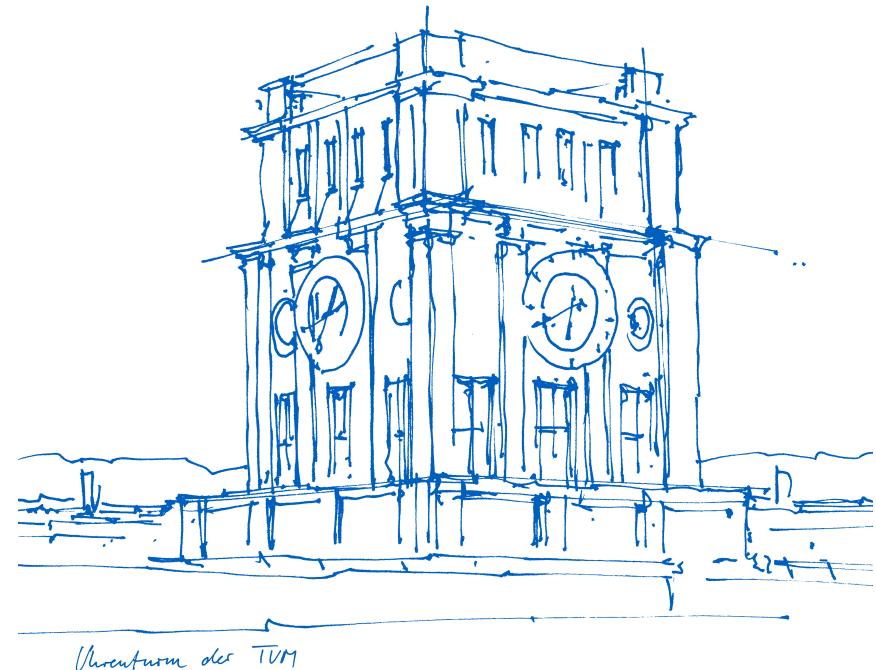
Martin Schulz

Exercises: Amir Raoofy

Technische Universität München

Fakultät für Informatik

Lecture 13:
Alternative Parallel
Programming Models



Summary From Last Time



Accelerators can help boost per node performance

- Dedicated hardware with specialized features
- Examples: FPGAs, Dataflow engines, GPUs

GPUs have evolved to massively parallel compute engines

- Hierarchical structure
- Inner elements are SIMD

Impact on programmability

- Accelerators rely on a host system
- Specialized language (extension) to cross-compile and control

CUDA targeting NVIDIA GPUs specifically

- Manual specification of data transfers through CUDA runtime
- Manual specification of thread mapping

OpenMP supports accelerators since OpenMP 4.0

- Allows specification of kernels on a device
- Pragmas to create the necessary data environment

Addendum: Asynchrony in CUDA

```
//Kernel
__global__
void daxpy(int n, double a, double *x, double *y) {
    int i=blockIdx.x*blockDim.x+threadIdx.x;
    if (i<n) y[i]=a*x[i]+y[i]
}
```

```
//Copy data to device
cudaMalloc(&d_x, nbytes); cudaMalloc(&d_y, nbytes);
cudaMemcpy(d_x,h_x,nbytes,cudaMemcpyHostToDevice);
cudaMemcpy(d_y,h_y,nbytes,cudaMemcpyHostToDevice);

// invoke DAXPY with 256 threads per Thread Block
int nblocks = (n+255)/256;

daxpy<<<nblocks, 256>>>(n,2.0,d_x,d_y);

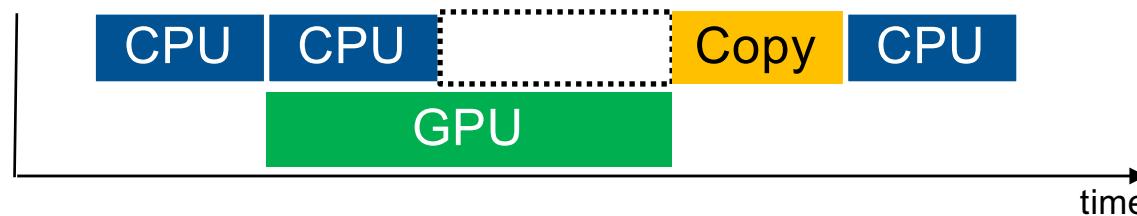
//Copy data from device
cudaMemcpy(h_y,d_y,nbytes,cudaMemcpyDeviceToHost);
```



Asynchrony Rules in CUDA

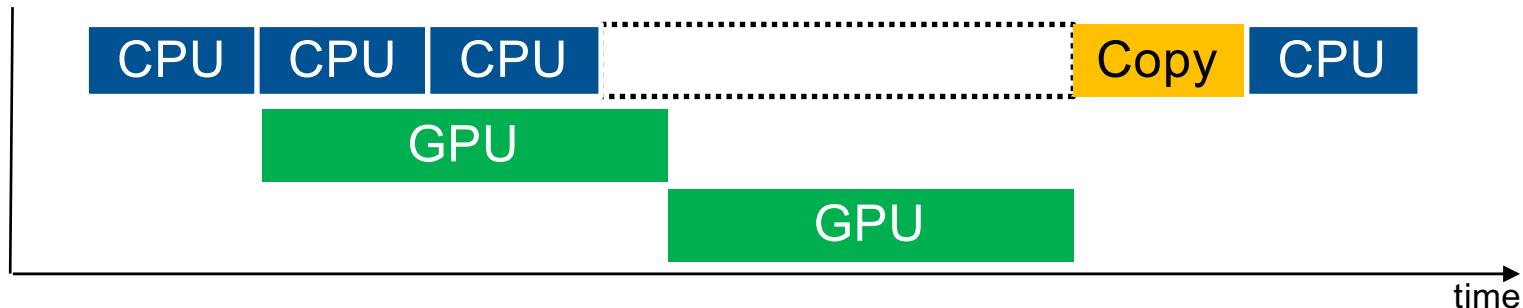
All CUDA kernel invocations are asynchronous

- Kernel just scheduled on GPU
- CPU continues executing
- Synchronization through explicit sync routines or synchronizing MemCopy



However, multiple kernel invocations serialize by default

- Each kernel waiting on the completion of the previous one
- Synchronizing MemCopy waiting on all kernels



CUDA Streams



Streams represents queues for kernel execution

- Host CPU code places work into the stream/queue and continues
- GPU schedules kernels from queue when resources are free

Overlapping

- Operations within the same stream are ordered (FIFO) and cannot overlap
- Operations in different streams are unordered and can overlap

Stream Management

`cudaStream_t stream;`

Declares a stream handle

`cudaStreamCreate(&stream);`

Allocates a stream

`cudaStreamDestroy(stream);`

Deallocates a stream

Synchronizes host until work in stream has completed

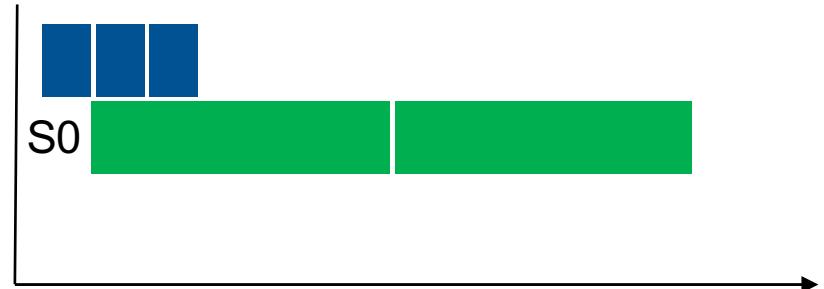
Kernel invocation on particular stream

`name<<blocks, threads, size, stream>>(<params>)`

CUDA Stream Example

Starting two kernels

```
foo<<blocks,threads>>();  
foo<<blocks,threads>>();
```



Starting two kernels on two streams

```
cudaStream_t stream1, stream2;  
cudaStreamCreate(&stream1);  
cudaStreamCreate(&stream2);
```

```
foo<<blocks,threads,0,stream1>>();  
foo<<blocks,threads,0,stream2>>();
```

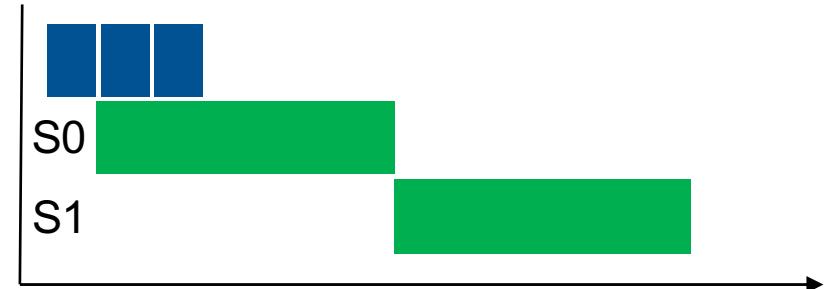
```
cudaStreamDestroy(stream1);  
cudaStreamDestroy(stream2);
```



Default Stream 0

Starting two kernels with one extra stream

```
cudaStream_t stream;  
cudaStreamCreate(&stream);
```

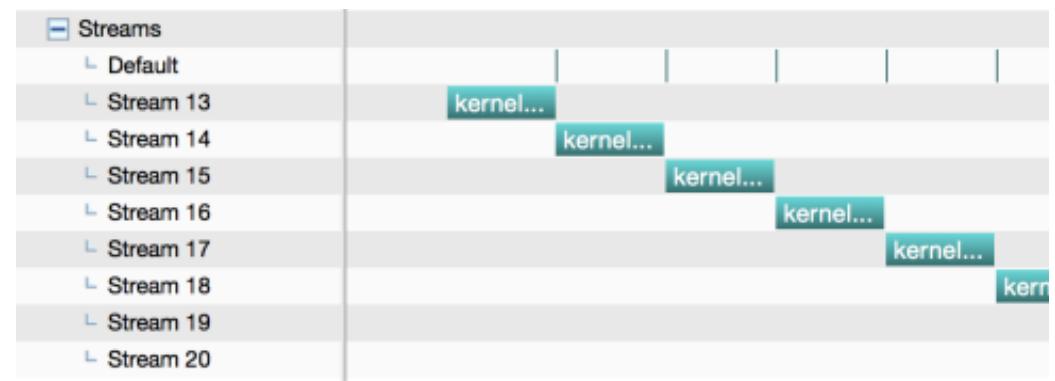


```
foo<<blocks,threads>>();  
foo<<blocks,threads,0,stream>>();  
  
cudaStreamDestroy(stream);
```

Stream 0 is the default stream

Stream is synchronizing to all streams

```
for (int i=0; i<ns; i++)  
{  
    cudaStreamCreate(&s[i]);  
    kernel<<<1,64,0,s[i]>>>();  
    kernel<<<1, 1>>>();  
}
```



Per Stream Memory Copies

Synchronizing MemCopy (synchronizes to stream 0)

```
cudaMemcpy (void* d, void* s, size_t c, cudaMemcpyKind kind)
```

Async MemCopy

```
cudaMemcpyAsync (void* d, void* s, size_t c,  
                 cudaMemcpyKind kind,  
                 cudaStream_t stream = 0 )
```

Async copy requires pinned host memory

- Memory not pageable, i.e., can be safely transferred at any time
- Special allocation/deallocation calls

```
cudaMallocHost(...) / cudaHostAlloc(...)  
cudaFreeHost(...)
```

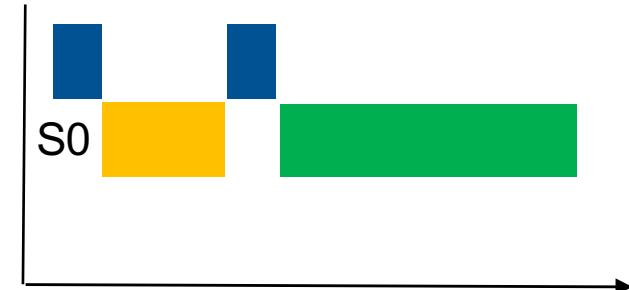
- Calls to pin/unpin memory

```
cudaHostRegister(...) / cudaHostUnregister(...)
```

Streams Can Also Help Transfer Concurrency

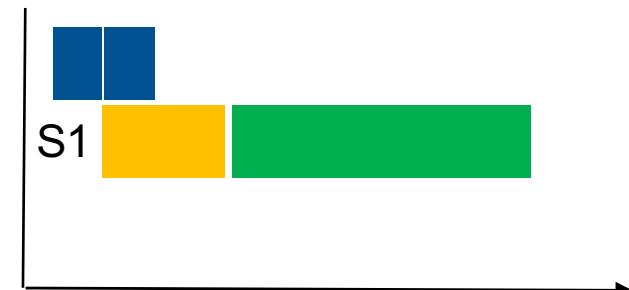
Synchronous

```
cudaMemcpy( ... );  
foo<<...>>();
```



Asynchronous Same Stream

```
cudaMemcpyAsync( ..., stream1 );  
foo<<..., stream1>>();
```



Asynchronous Different Streams

```
cudaMemcpyAsync( ..., stream1 );  
foo<<..., stream2>>();
```



Advanced CUDA Synchronization Topics

Explicit Synchronization

- Synchronize everything: `cudaDeviceSynchronize()`
- Synchronize one stream: `cudaStreamSynchronize(stream)`
- Synchronize using events -> look it up

Stream Priorities

- High priority streams will preempt lower priority streams
 - `cudaDeviceGetStreamPriorityRange(&low, &high)`
 - Lower number is higher priority
- Create using special API
 - `cudaStreamCreateWithPriority(&stream, flags, priority)`

Stream Callbacks

- Registration of host side callbacks with respect to a stream
- Invocation when tasks in the stream are complete

More information

- <http://on-demand.gputechconf.com/gtc/2014/presentations/S4158-cuda-streams-best-practices-common-pitfalls.pdf>
- <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency/>

Programming Models/Approaches So Far

Several on-node models

- Pthreads
 - Explicit access to hardware threads and synchronization
- OpenMP
 - Loop and task parallelism with more structured synchronization
- CUDA
 - NVIDIA GPU programming
- Focus on shared address spaces (even CUDA to some degree)

Only cross-node model: MPI

- Explicit message passing
- Point-to-Point communication as well as collectives
- Newer versions provide one-sided communication
 - Limited version of a shared address space

Question:

Can a global address space also work across nodes?

Partitioned Global Address Space

Global address space

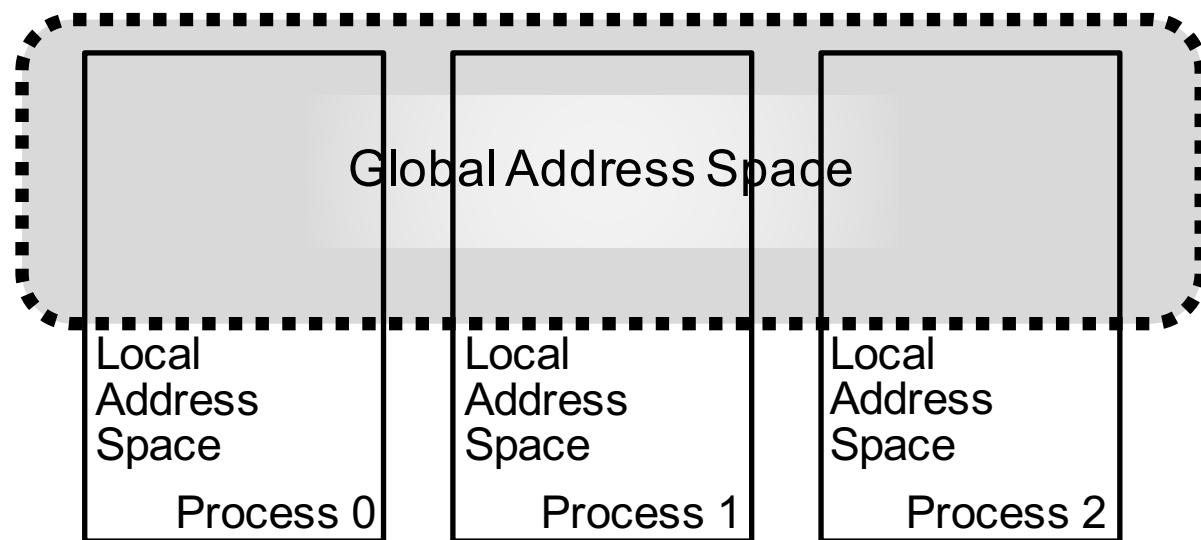
- Any thread/process may directly read/write data allocated by another

Partitioned:

- Data is designated as local or global

By default:

- Heaps are Global
- Stacks are Local
- Can be adjusted



Examples

- SPMD languages: UPC, CAF, and Titanium
- All three use an SPMD execution model

Dynamic languages: X10, Fortress, Chapel and Charm++

Example: UPC

Unified Parallel C [LBL, UC Berkeley, GWU]

Both private and shared data

`int x[10];` vs. `shared int y[10];`

One-sided shared-memory communication through simple assignment:

`x[i] = y[i];` or `t = *p;`

Support for distributed data structures

- Distributed arrays; local and global pointers/references

Synchronization

- Global barriers, locks, memory fences

Collective Communication, IO libraries, etc.

UPC Example

```
#define SIZE 4*THREADS
shared int A[SIZE];

int main()
{
    int i;
    upc_forall(i=0; i < SIZE; i++; &A[i])
    {
        A[i] = MYTHREAD;
    }
    upc_barrier;
    if (MYTHREAD == 0)
    {
        for (i=0; i < SIZE; i++)
            { printf("%d ", A[i]); }
        printf("\n");
    }
    return 0;
}
```

Affinity Expression

Local & Remote Reads from A

PGAS Pros & Cons

Advantages of PGAS

- Low latency memory access operations instead of two sided messaging
- Direct mapping of global data structures

Disadvantages of PGAS

- All the baggage of shared memory, incl. synchronization and possible races
- Harder to optimize
- Global memory accesses have larger latencies
- Only two hierarchy levels

Common with the other approaches (MPI, OpenMP, Pthreads, CUDA, ...)

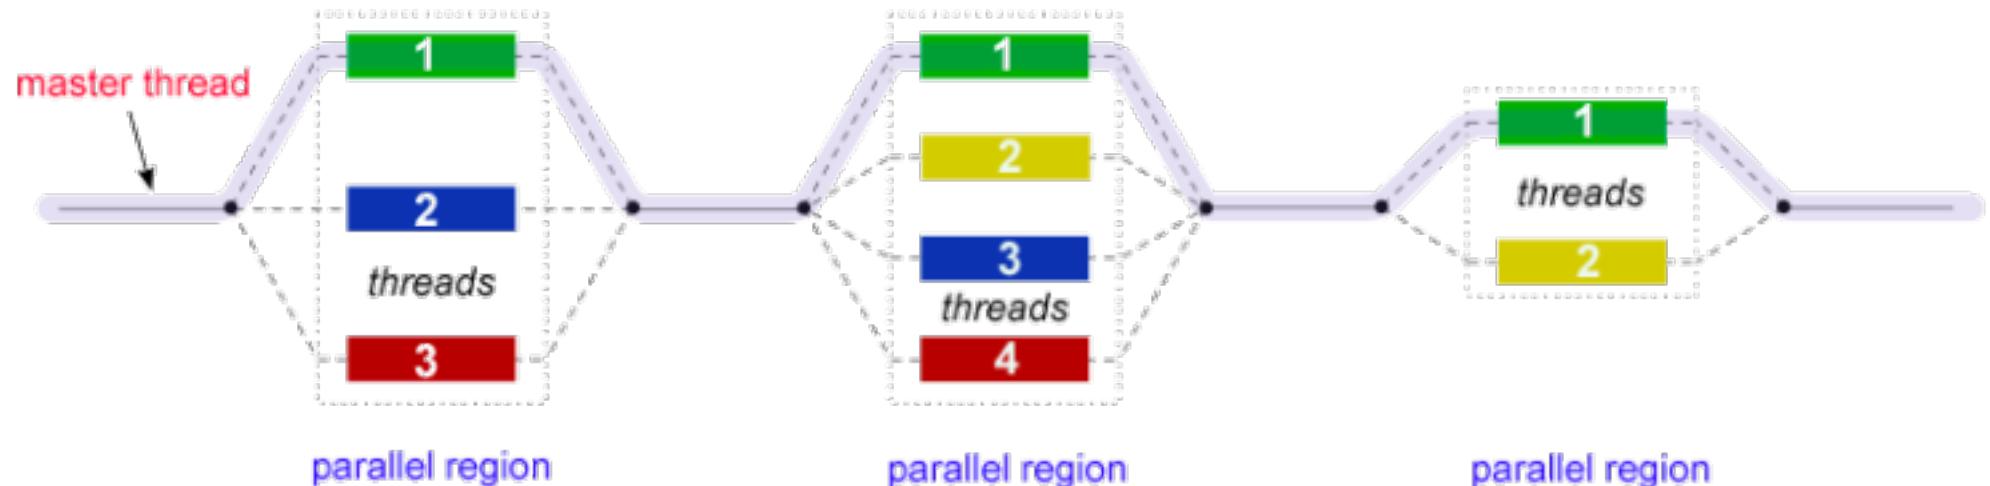
- Low-level and close to hardware
- Good abstractions for performance oriented programming
- Not well suited for user-level and application-level abstractions

Question:

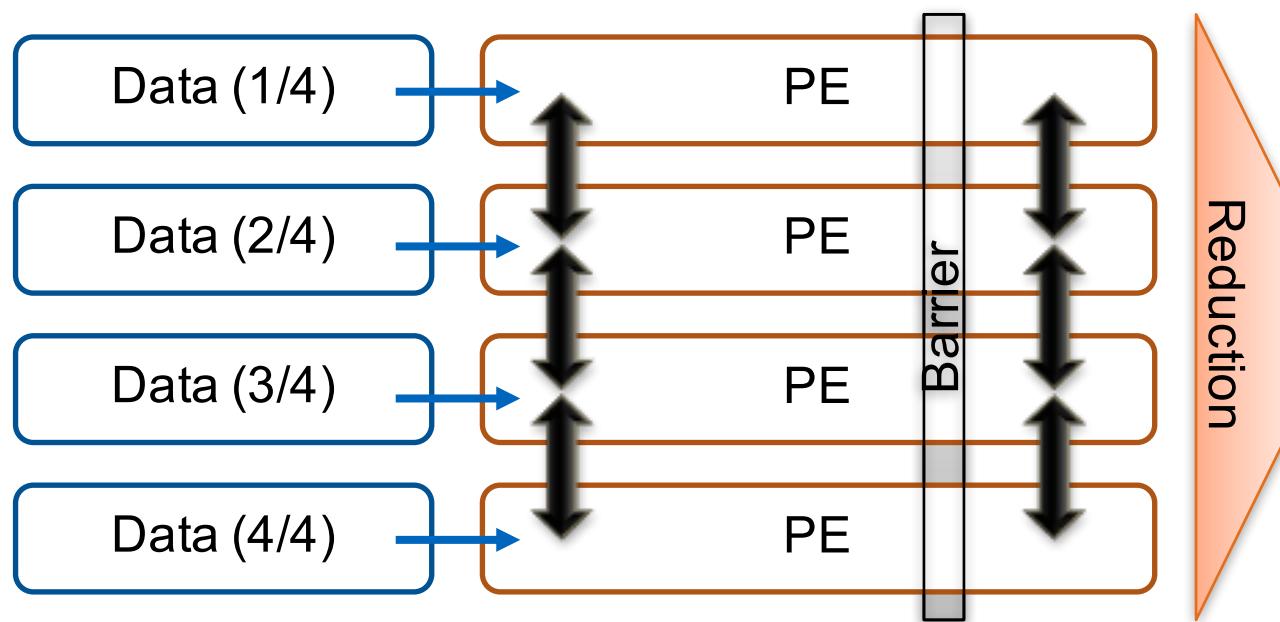
Can we find abstractions that are closer to the applications and its requirements?
How to get away from restrictions of SPMD and/or Fork/Join?

Traditional Parallelization Patterns

Fork/Join



SPMD



Task Graphs

Idea composition of the problem into basic elements

- Individual tasks
- Dependencies between tasks

Only have to fulfill dependencies

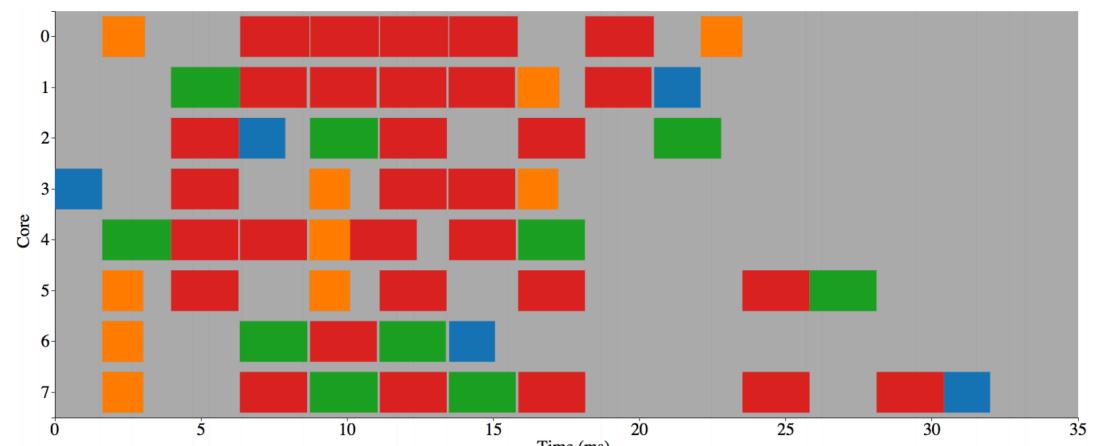
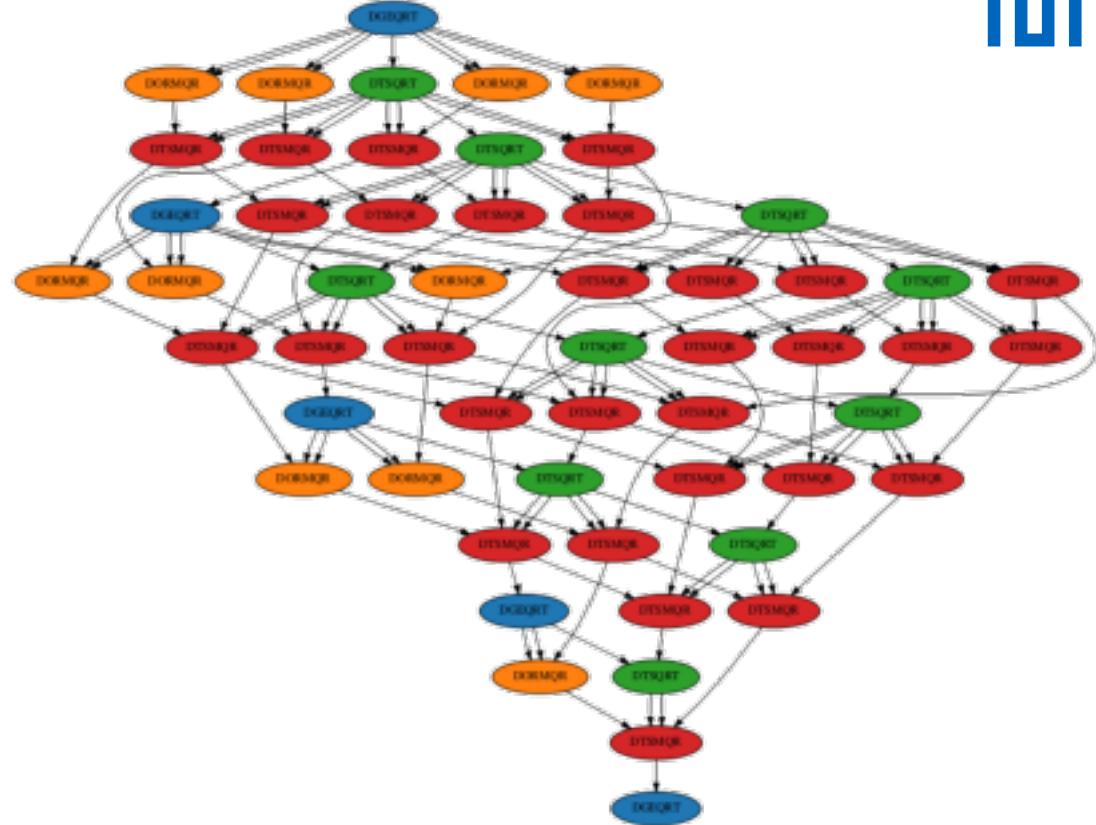
- No extra synchronization
- Scheduling freedom

Difficulties

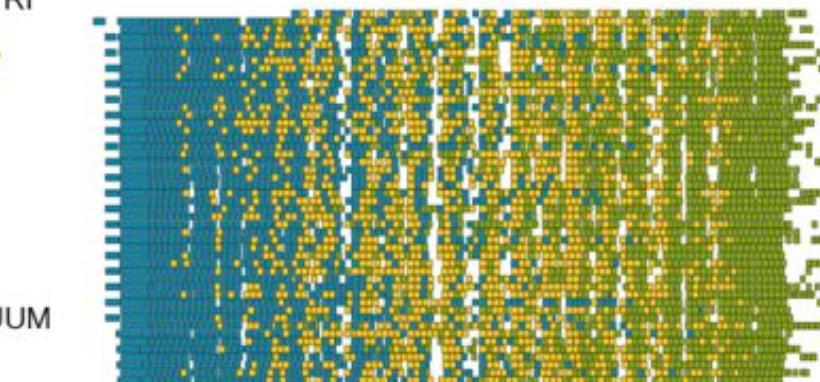
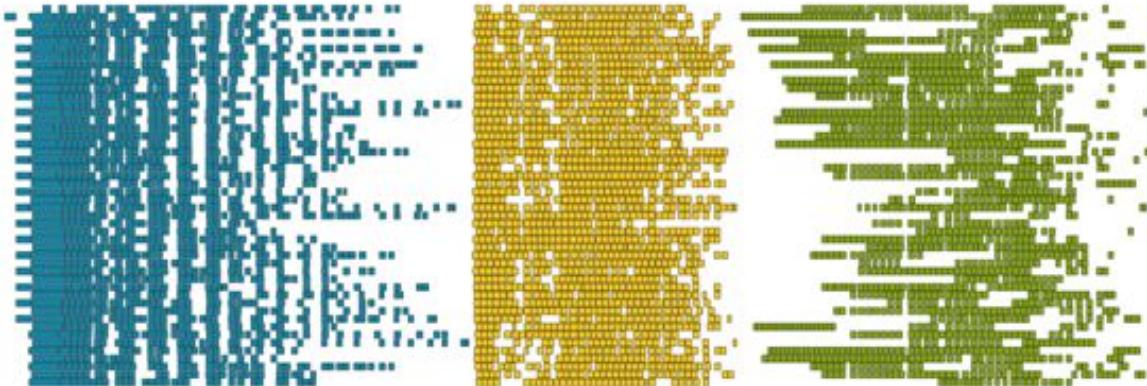
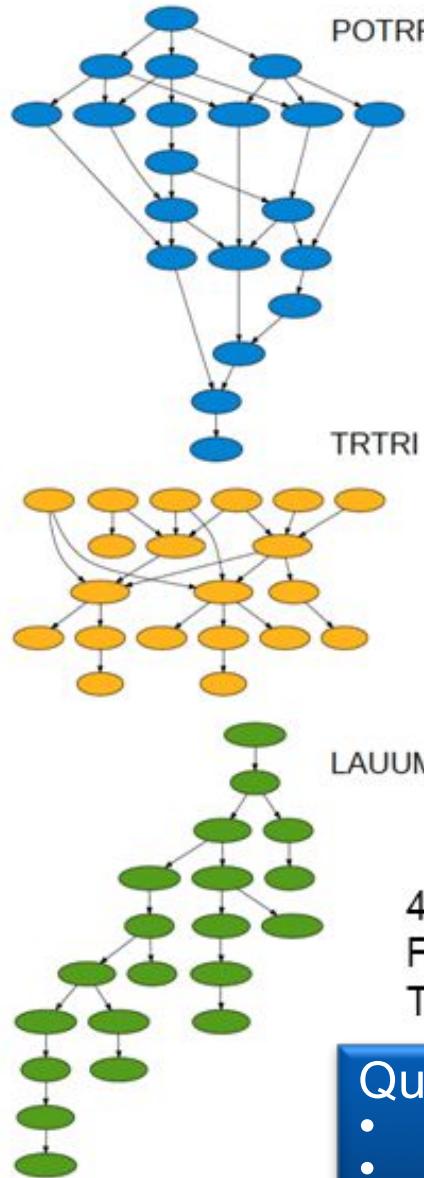
- Granularity
- Overhead

Challenge: schedule tasks

- Efficient mapping of tasks
- NUMA awareness
- Fast handover

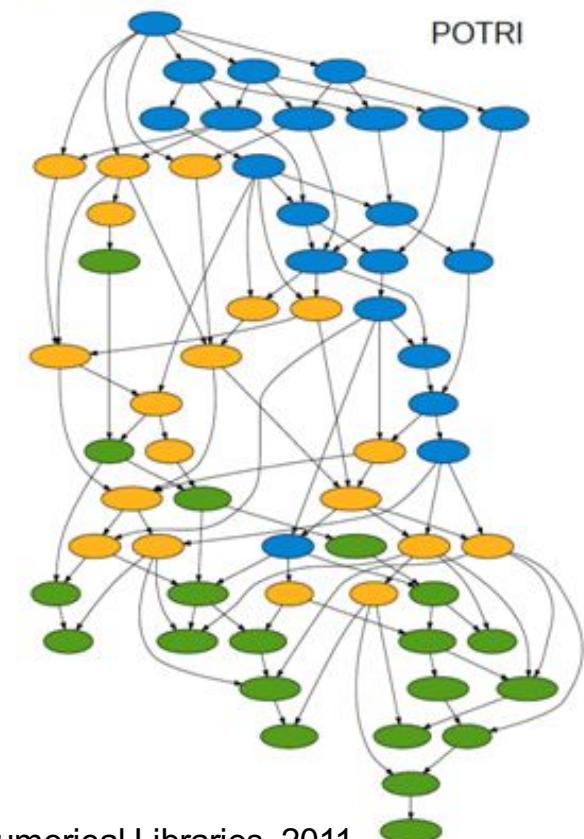


Pipelining: Cholesky Inversion



48 cores
POTRF, TRTRI and LAUUM.
The matrix is 4000 x 4000, tile size is 200 x 200,

Quark System at UTK
• Fine grain tasks
• Direct composition of numerical kernels



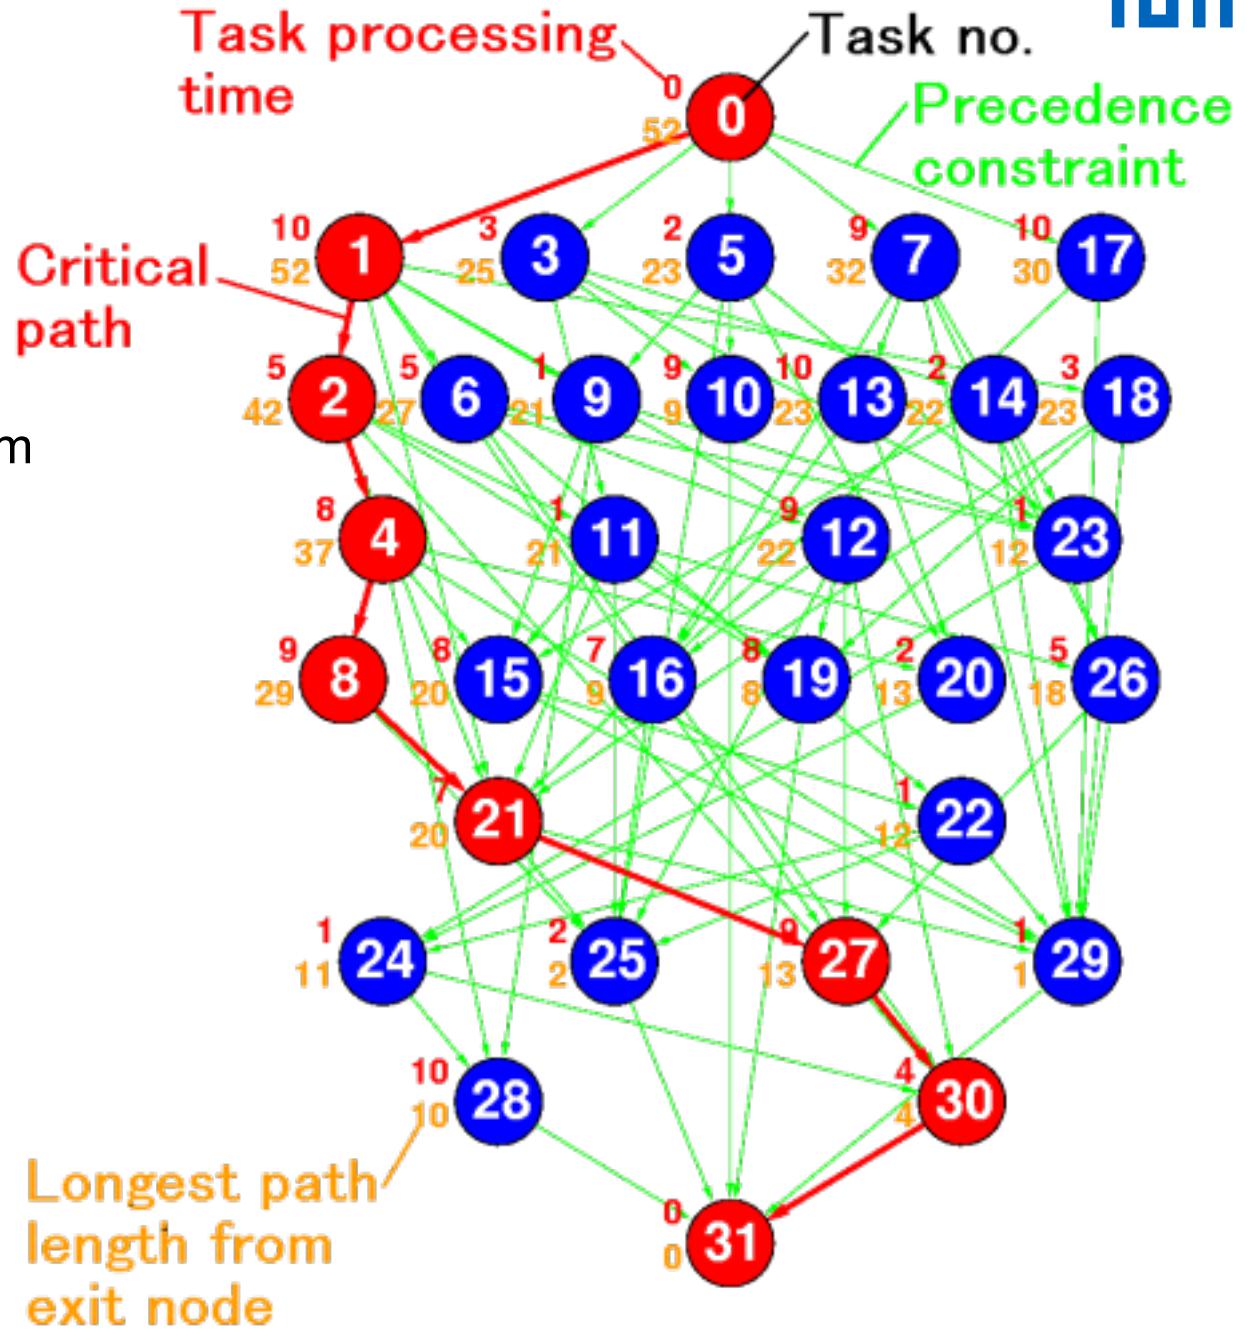
Critical Paths

Only relevant metric

- Length of critical path
- Longest dependency list between tasks in system

Consequences

- Dictates performance
- Focus performance optimization on this path



Two Task-Based Examples

Charm++

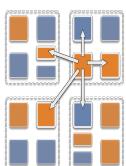
- Actor-Based programming language
- Developed at UIUC
- Slides by Prof. Dr. Laxmikant (Sanjay) Kale

Legion

- Data centric tasking runtime and programming model
- Developed at Stanford together with LANL and NVIDIA
- Slides by Prof. Dr. Alex Aiken

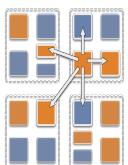
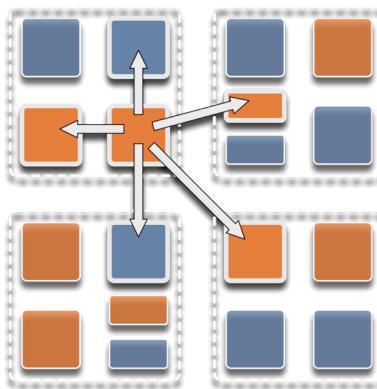
What is Charm++?

- Charm++ is a generalized approach to writing parallel programs
 - An alternative to the likes of MPI, UPC, GA etc.
 - But not to sequential languages such as C, C++, and Fortran
- Represents:
 - The style of writing parallel programs
 - The runtime system
 - And the entire ecosystem that surrounds it
- Three design principles:
 - Overdecomposition, Migratability, Asynchrony



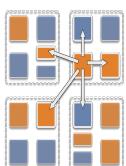
Overdecomposition

- Decompose the work units & data units into many more pieces than execution units
 - Cores/Nodes/...
- Not so hard: we do decomposition anyway



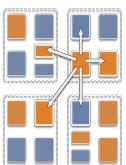
Migratability

- Allow these work and data units to be migratable at runtime
 - i.e. the programmer or runtime can move them
- Consequences for the application developer
 - Communication must now be addressed to logical units with global names, not to physical processors
 - But this is a good thing
- Consequences for RTS
 - Must keep track of where each unit is
 - Naming and location management

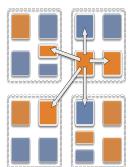
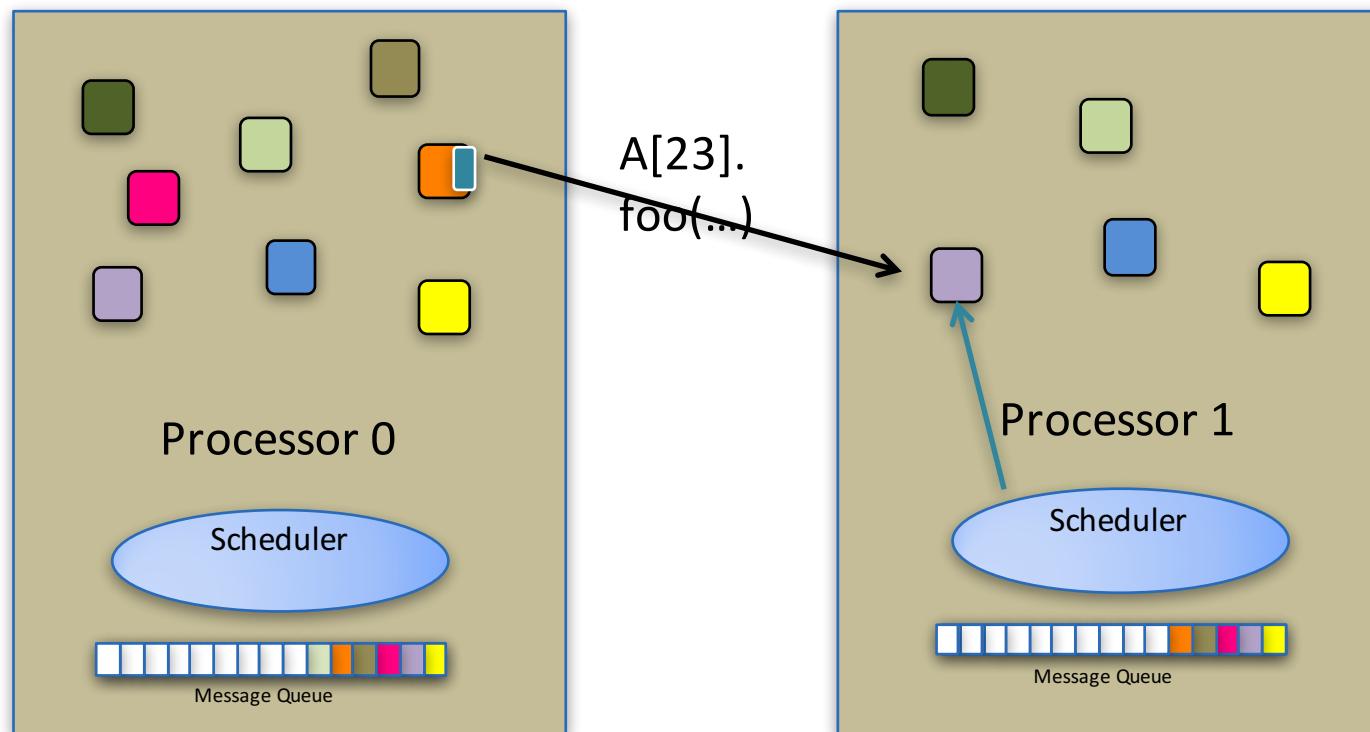


Asynchrony: Message-Driven Execution

- With over-decomposition and migratability:
 - You have multiple units on each processor
 - They address each other via logical names
- Need for scheduling:
 - What sequence should the work units execute in?
 - One answer: let the programmer sequence them
 - Seen in current codes, e.g. some AMR frameworks
 - Message-driven execution:
 - Let the work-unit that happens to have data (“message”) available for it execute next
 - Let the RTS select among ready work units
 - Programmer should not specify what executes next, but can influence it via priorities

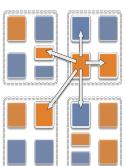


Message-driven Execution



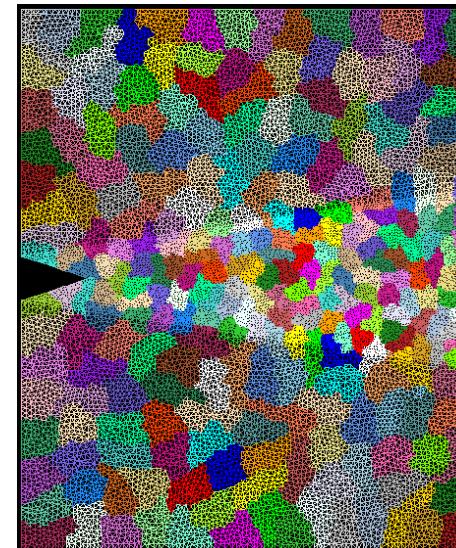
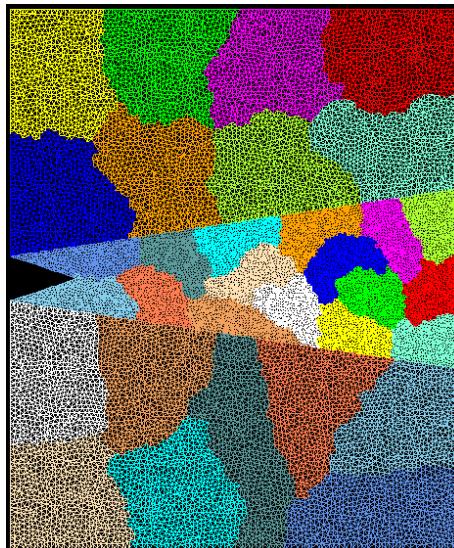
Grainsize

- Charm++ philosophy:
 - Let the programmer decompose their work and data into coarse-grained entities
- It is important to understand what I mean by coarse-grained entities
 - You don't write sequential programs that some system will auto-decompose
 - You don't write programs when there is one object for each *float*
 - You consciously choose a grainsize, but choose it independently of the number of processors
 - Or parameterize it, so you can tune later

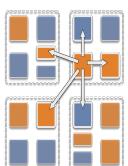


Crack Propagation

This is 2D, circa 2002...
but shows overdecomposition for unstructured meshes

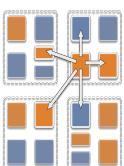


Decomposition into 16 chunks (left) and 128 chunks, 8 for each PE (right).
The middle area contains cohesive elements. Both decompositions obtained
using Metis. Pictures: S. Breitenfeld, and P. Geubelle



Realization of This Model in Charm++

- Overdecomposed entities: chares
 - Chares are C++ objects
 - With methods designated as “entry” methods
 - Which can be invoked asynchronously by remote chares
 - Chares are organized into indexed collections
 - Each collection may have its own indexing scheme
 - 1D, ..., 6D
 - Sparse
 - Bitvector or string as an index
 - Chares communicate via asynchronous method invocations
 - `A[i].foo(...);`
 - A is the name of a collection, i is the index of the particular chare.



Empowering the RTS

Adaptive
Runtime System

Introspection

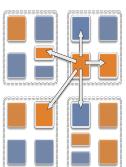
Adaptivity

Asynchrony

Overdecomposition

Migratability

- The Adaptive RTS can:
 - Dynamically balance loads
 - Optimize communication:
 - Spread over time, async collectives
 - Automatic latency tolerance
 - Prefetch data with almost perfect predictability



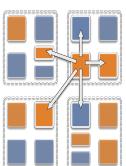
Hello World Example

- hello.ci file

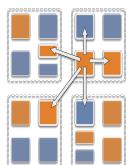
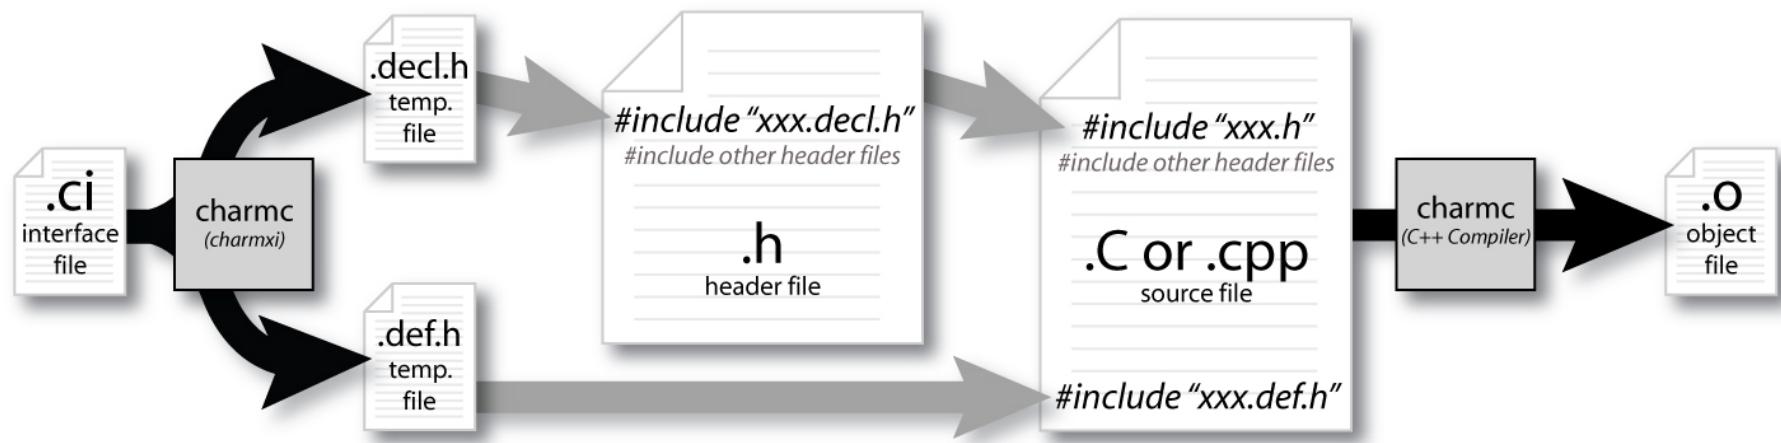
```
mainmodule hello {  
    mainchare Main {  
        entry Main(CkArgMsg *m);  
    };  
};
```

- hello.cpp file

```
#include <stdio.h>  
#include "hello.decl.h"  
  
class Main : public CBase_Main {  
public: Main(CkArgMsg* m) {  
    ckout << "Hello World!" << endl;  
    CkExit();  
};  
};  
  
#include "hello.def.h"
```



Compiling a Charm++ Program

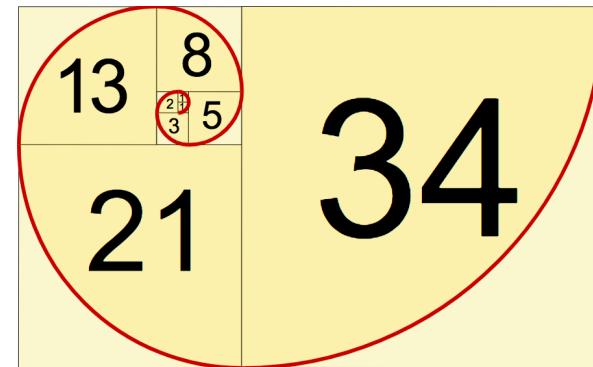


Example: Fibonacci (.ci file)

```
mainmodule fib {
    readonly int n;

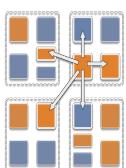
    mainchare Main {
        entry Main(CkArgMsg* m);
        entry void done();
    };

    chare Fib {
        entry Fib(int n, bool is_root, CProxy_Fib parent);
        entry void calc(int n) {
            if (n < THRESHOLD)
                serial { respond(seqFib(n)); }
            else {
                serial {
                    CProxy_Fib::ckNew(n - 1, false, thisProxy);
                    CProxy_Fib::ckNew(n - 2, false, thisProxy);
                }
                when response(int val)
                    when response(int val2)
                        serial { respond(val + val2); }
                }
            };
            entry void response(int val);
        };
    };
}
```



From: mathisfun.com

$$F(n) = F(n-1) + F(n-2)$$



Example: Fibonacci (.cpp file, part 1)

```
#include "fib.decl.h"

#define THRESHOLD 3 // calculate sequentially if below

/* readonly */ int n; // fibonacci number
/* readonly */ CProxy_Main mainProxy;

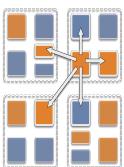
class Main : public CBase_Main {
    double start_time;

public:
    Main(CkArgMsg* m) {
        // determine fibonacci number
        n = 20;
        if (m->argc == 2)
            n = atoi(m->argv[1]);

        CkPrintf("\n[Exercise 1. Fibonacci Sequence]\n");

        // create chores and start computation
        start_time = CkWallTimer();
        CProxy_Fib::ckNew(n, true, CProxy_Fib());
    }

    void done() {
        CkPrintf("Elapsed time: %lf seconds\n\n", CkWallTimer()-start_time);
        CkExit();
    }
}
```



Example: Fibonacci (.cpp file, part 2)

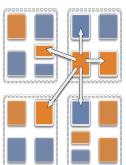
```
class Fib : public CBase_Fib {
    Fib_SDAG_CODE
    CProxy_Fib parent;
    bool is_root;

public:
    Fib(int n, bool is_root_, CProxy_Fib parent_)
        : parent(parent_), is_root(is_root_) {
        thisProxy.calc(n);
    }

    int seqFib(int n) {
        return (n < 2) ? n : seqFib(n - 1) + seqFib(n - 2);
    }

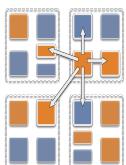
    void respond(int val) {
        if (!is_root) {
            parent.response(val);
            delete this;
        }
        else {
            CkPrintf("Fibonacci number is: %d\n", val);
            mainProxy.done();
        }
    }
};

#include "fib.def.h"
```



Summary: What is Charm++?

- Charm++ is a way of parallel programming
- It is based on:
 - Objects
 - Overdecomposition
 - Asynchrony
 - Asynchronous method invocations
 - Migratability
 - Adaptive runtime system
- It has been co-developed synergistically with multiple CSE applications



Two Task-Based Examples

Charm++

- Actor-Based programming language
- Developed at UIUC
- Slides by Prof. Dr. Laxmikant (Sanjay) Kale

Legion

- Data centric tasking runtime and programming model
- Developed at Stanford together with LANL and NVIDIA
- Slides by Prof. Dr. Alex Aiken



Legion Approach

- Asynchronous tasking

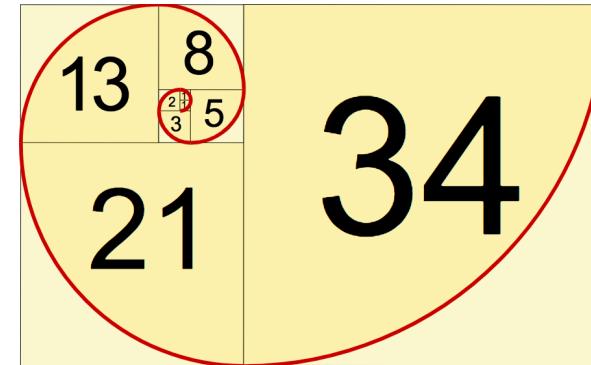


Fibonacci (definitions)

```
#include <cstdio>
#include <cassert>
#include <cstdlib>
#include "legion.h"

using namespace Legion;

enum TaskIDs
{
    TOP_LEVEL_TASK_ID,
    FIBONACCI_TASK_ID,
    SUM_TASK_ID,
};
```



From: mathisfun.com

$$\begin{aligned} F(n) = \\ F(n-1) + \\ F(n-2) \end{aligned}$$



Fibonacci (main)

```
int main(int argc, char **argv)
{
    Runtime::set_top_level_task_id(TOP_LEVEL_TASK_ID);
    {
        TaskVariantRegistrar registrar(TOP_LEVEL_TASK_ID, "top_level");
        registrar.add_constraint(ProcessorConstraint(Processor::LOC_PROC));
        Runtime::preregister_task_variant<top_level_task>(registrar, "top_level");
    }
    {
        TaskVariantRegistrar registrar(FIBONACCI_TASK_ID, "fibonacci");
        registrar.add_constraint(ProcessorConstraint(Processor::LOC_PROC));
        Runtime::preregister_task_variant<int, fibonacci_task>(registrar, "fibonacci");
    }
    {
        TaskVariantRegistrar registrar(SUM_TASK_ID, "sum");
        registrar.add_constraint(ProcessorConstraint(Processor::LOC_PROC));
        registrar.set_leaf(true);
        Runtime::preregister_task_variant<int, sum_task>(registrar, "sum",
                                                       AUTO_GENERATE_ID);
    }

    return Runtime::start(argc, argv);
}
```



Fibonacci (top-level task)

```
void top_level_task(const Task *task,
                     const std::vector<PhysicalRegion> &regions,
                     Context ctx, Runtime *runtime)
{
    int num_fibonacci = 7; // Default value
    const InputArgs &command_args = Runtime::get_input_args();

    for (int i = 1; i < command_args(argc; i++)
    {
        // Skip any legion runtime configuration parameters
        if (command_args.argv[i][0] == '-') { i++; continue; }
        num_fibonacci = atoi(command_args.argv[i]);
        break;
    }
    printf("Computing the first %d Fibonacci numbers...\n", num_fibonacci);
    std::vector<Future> fib_results;
    for (int i = 0; i < num_fibonacci; i++)
    {
        TaskLauncher launcher(FIBONACCI_TASK_ID, TaskArgument(&i,sizeof(i)));
        fib_results.push_back(runtime->execute_task(ctx, launcher));
    }
    for (int i = 0; i < num_fibonacci; i++)
    {
        int result = fib_results[i].get_result<int>();
        printf("Fibonacci(%d) = %d\n", i, result);
    }
    fib_results.clear();
}
```



Fibonacci (fib task)

```
int fibonacci_task(const Task *task,
                    const std::vector<PhysicalRegion> &regions,
                    Context ctx, Runtime *runtime)
{
    assert(task->arglen == sizeof(int));
    int fib_num = *(const int*)task->args;

    if (fib_num == 0) return 0;
    if (fib_num == 1) return 1;

    // Launch fib-1
    const int fib1 = fib_num-1;
    TaskLauncher t1(FIBONACCI_TASK_ID, TaskArgument(&fib1,sizeof(fib1)));
    Future f1 = runtime->execute_task(ctx, t1);

    // Launch fib-2
    const int fib2 = fib_num-2;
    TaskLauncher t2(FIBONACCI_TASK_ID, TaskArgument(&fib2,sizeof(fib2)));
    Future f2 = runtime->execute_task(ctx, t2);

    TaskLauncher sum(SUM_TASK_ID, TaskArgument(NULL, 0));
    sum.add_future(f1);
    sum.add_future(f2);

    Future result = runtime->execute_task(ctx, sum);
    return result.get_result<int>();
}
```



Fibonacci (sum task)

```
int sum_task(const Task *task,
             const std::vector<PhysicalRegion> &regions,
             Context ctx, Runtime *runtime)
{
    assert(task->futures.size() == 2);

    Future f1 = task->futures[0];
    int r1 = f1.get_result<int>();

    Future f2 = task->futures[1];
    int r2 = f2.get_result<int>();

    return (r1 + r2);
}
```

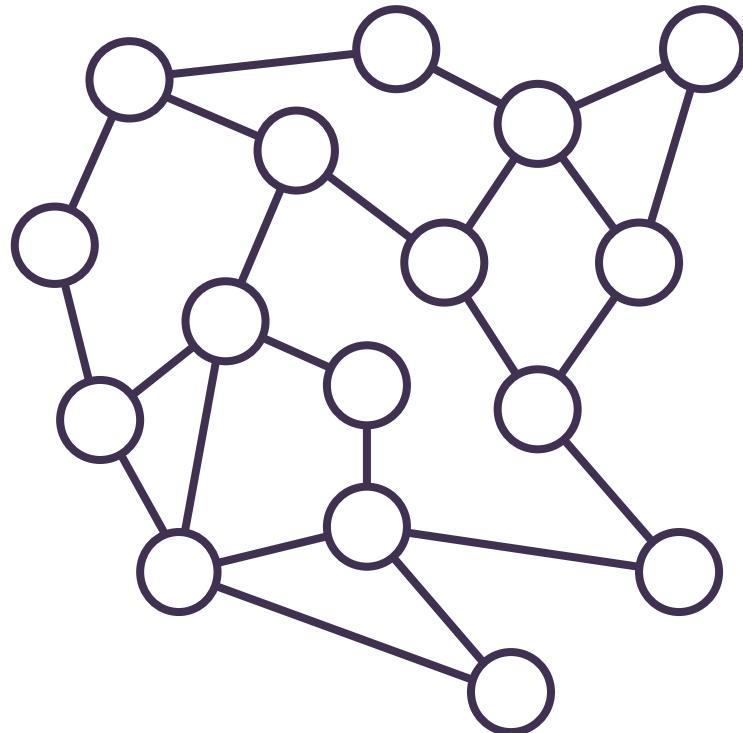


Legion Approach

- Asynchronous tasking
- Capture the structure of program data
- Decouple specification from *mapping*
- Automate
 - data movement
 - parallelism discovery
 - synchronization
 - hiding long latency operations



Example: Circuit Simulation





Example: Circuit Simulation

```
task simulate_circuit(Region[Node] N, Region[Wires] W)
```

```
{
```

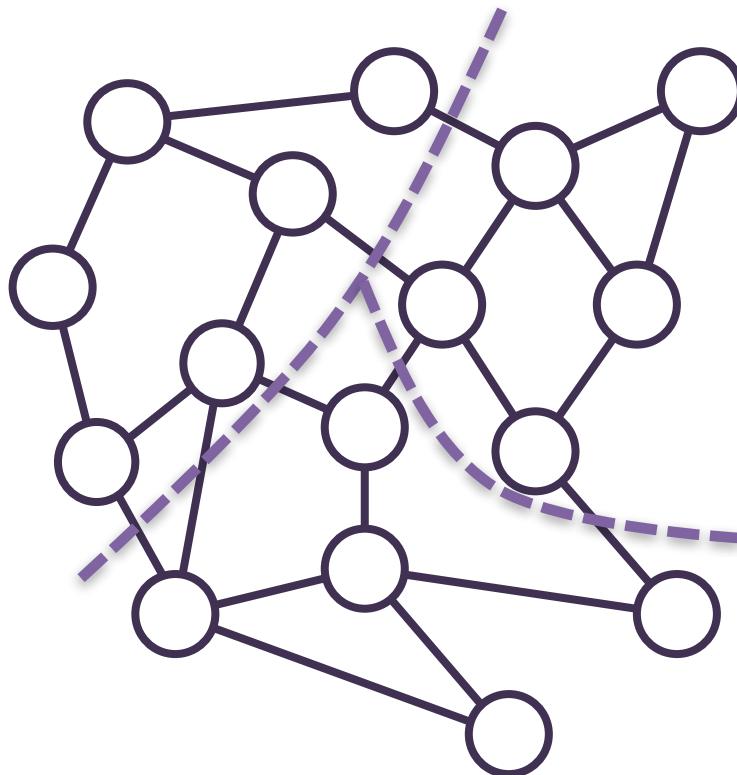
Tasks are the unit of parallel execution.

```
}
```

Logical regions are (typed) collections
Logical:
 no implied layout
 no implied location



Partitioning



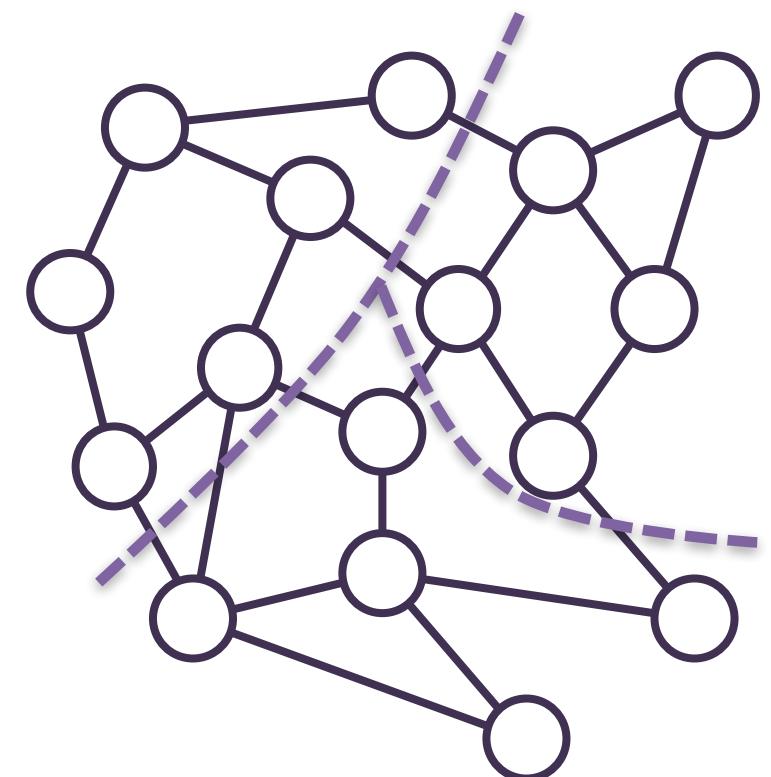
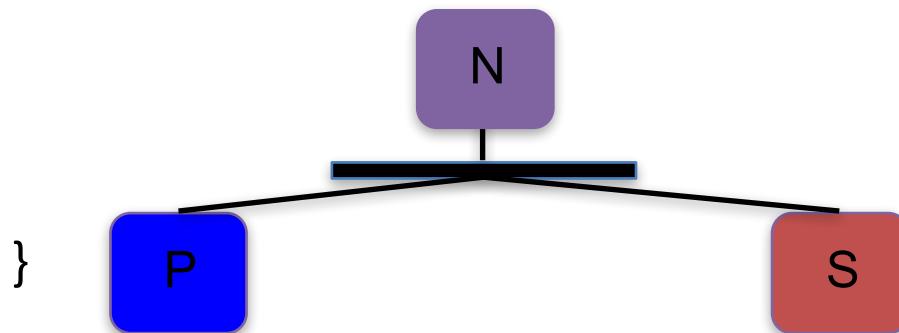


Partitioning

```
task simulate_circuit(Region[Node] N, Region[Wires] W)
```

```
{
```

```
[ P, S ] = partition(ps_map, N)
```

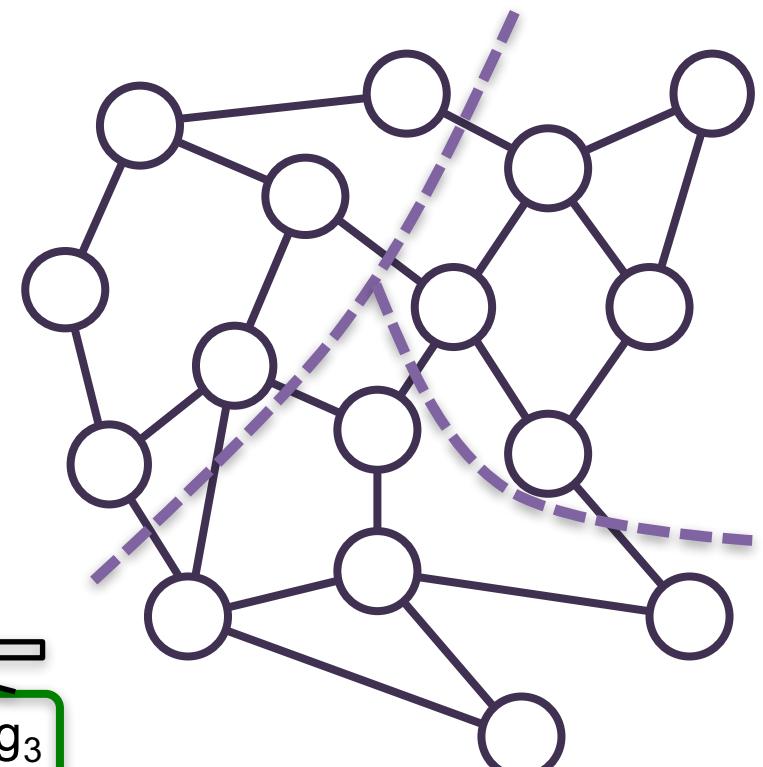
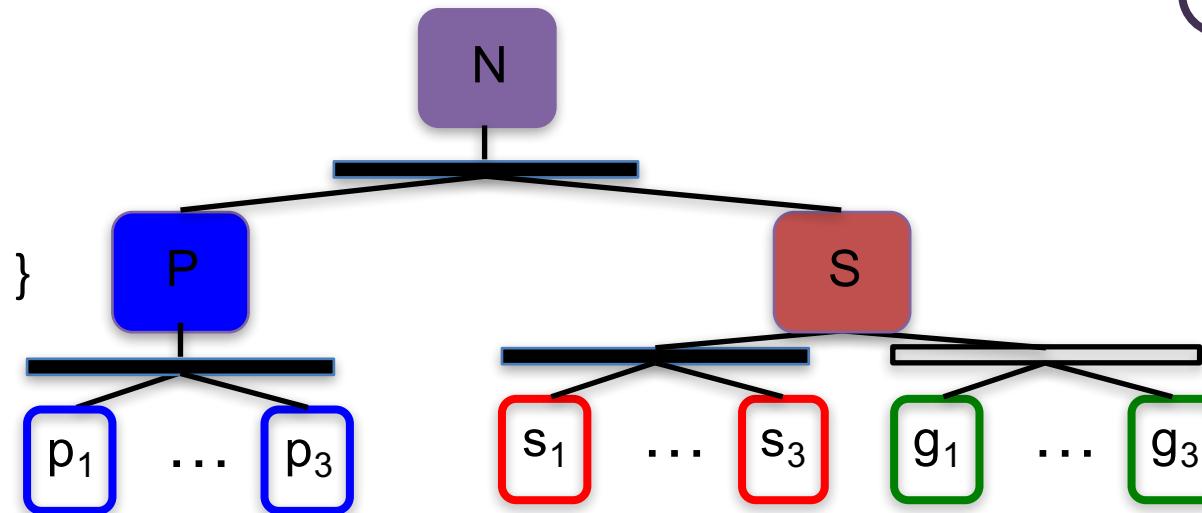




Partitioning

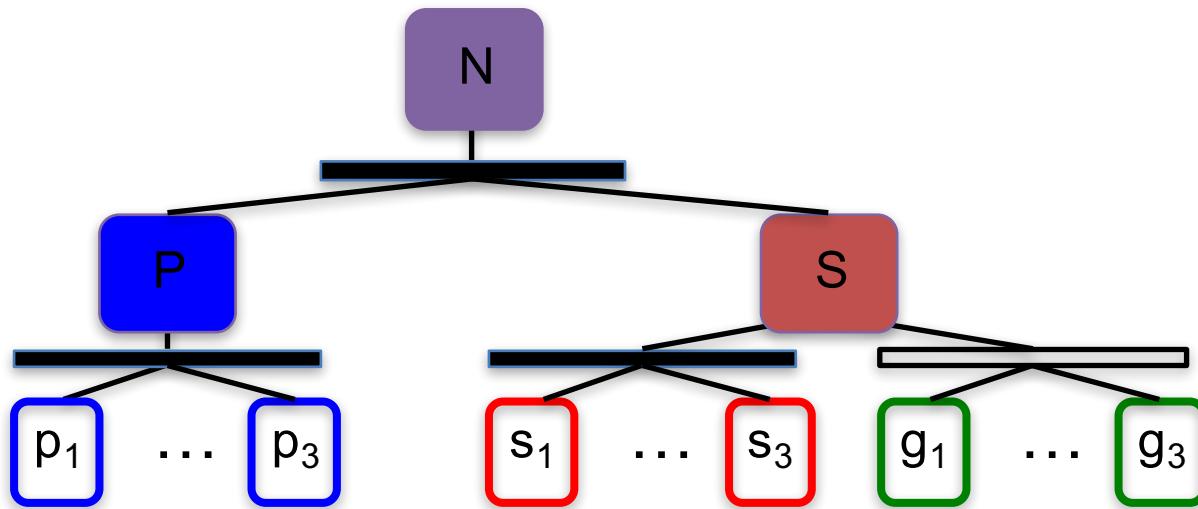
```
task simulate_circuit(Region[Node] N, Region[Wires] W)
```

```
{  
    Array[Region[Node]] private, shared, ghost;  
    ...  
    [ P, S ] = partition(ps_map, N)  
    private = partition(private_map, P)  
    shared = partition(shared_map, S)  
    ghost = partition(ghost_map, S)
```



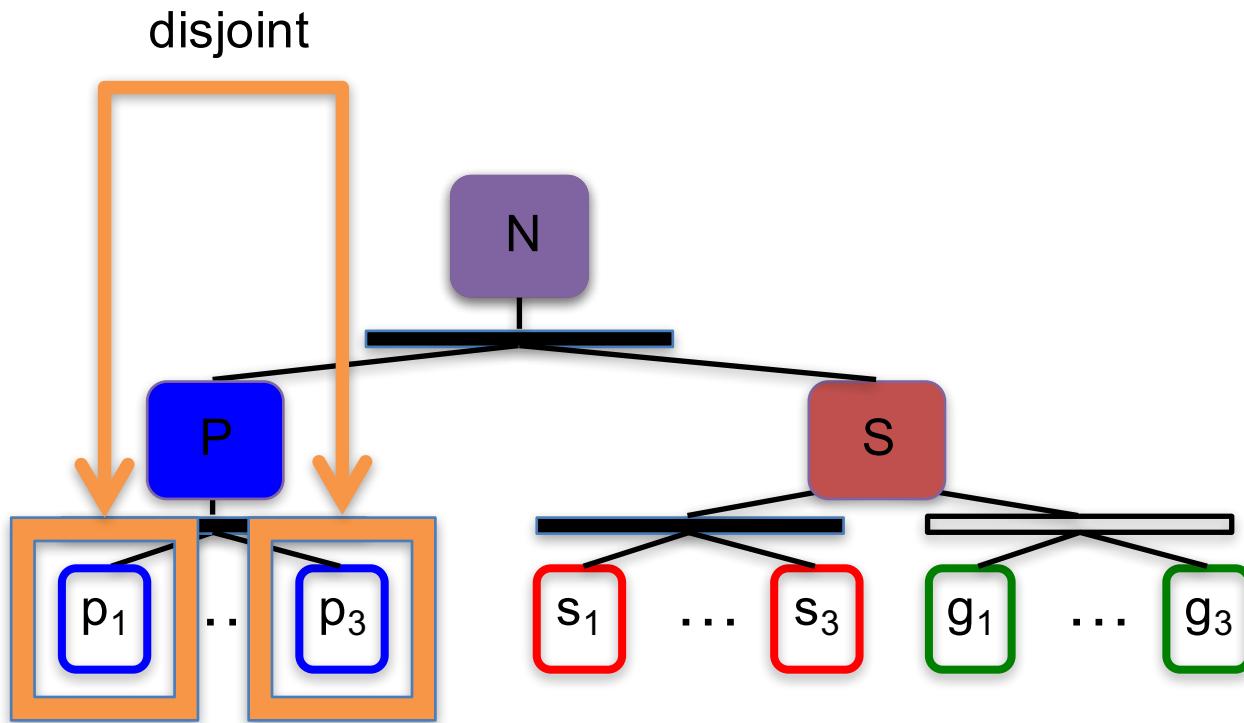


Region Trees





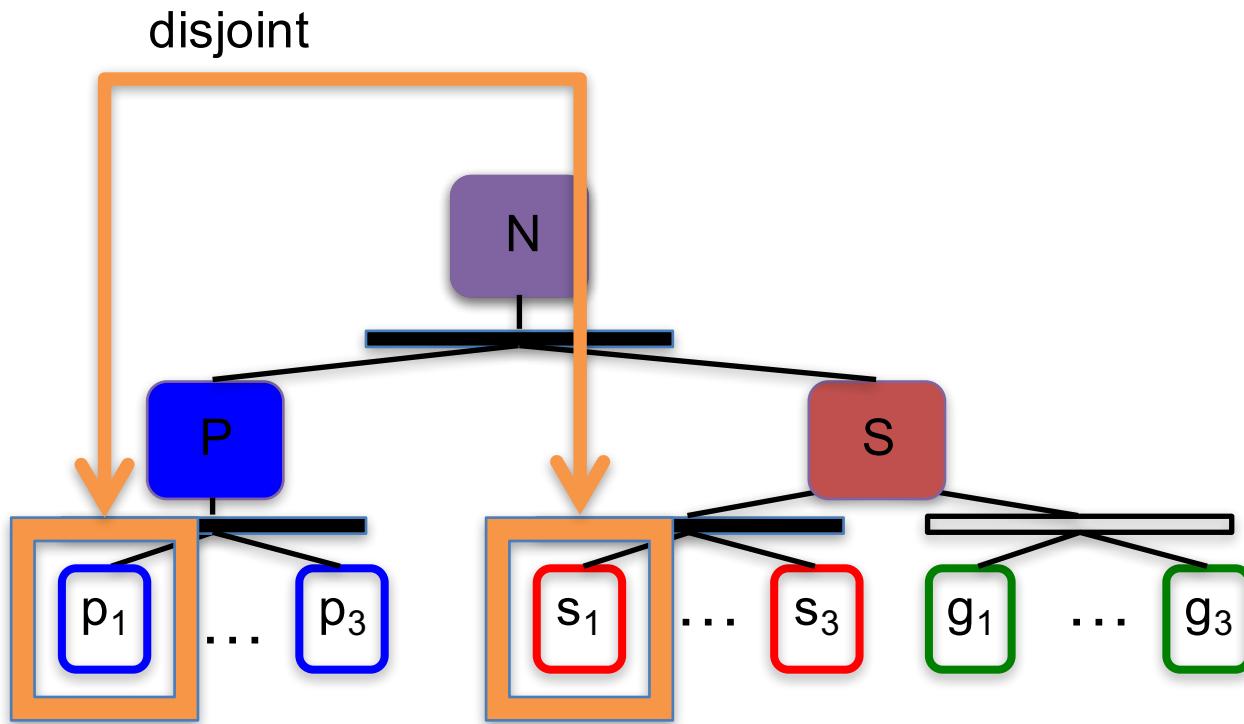
Region Trees



Locality
Independence



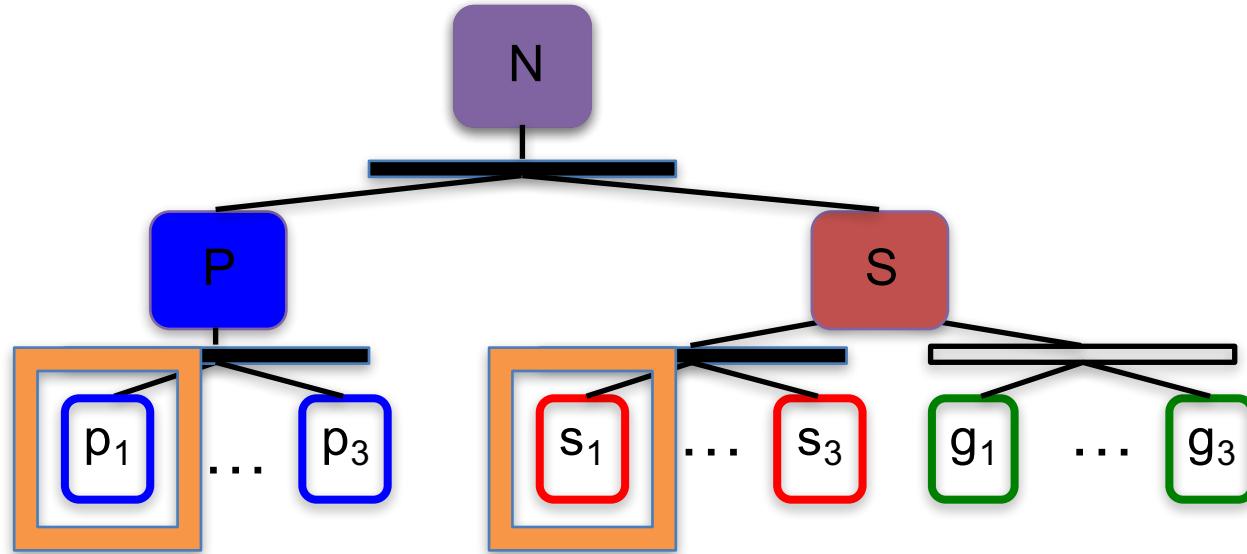
Region Trees



Locality
Independence



Region Trees

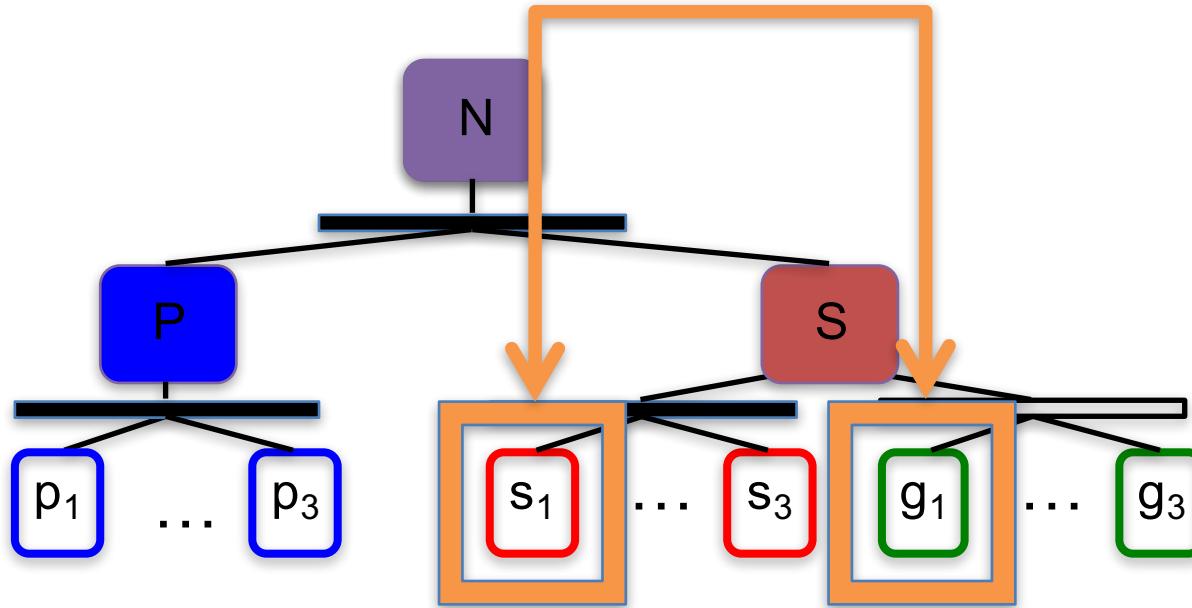


Locality
Independence
Aliasing



Region Trees

possibly overlap



Locality
Independence
Aliasing



Legion Tasks & Privileges

```
task simulate_circuit(Region[Node] N, Region[Wires] W) :  
    ReadWrite(N, W)
```

```
{
```

```
...
```

```
calc_currents(piece[0], p0, s0, g0);  
calc_currents(piece[1], p1, s1, g1);  
distribute_charge(piece[0], p0, s0, g0);  
distribute_charge(piece[1], p1, s1, g1);  
...
```

```
}
```

```
task calc_currents(Piece p) :  
    Read(p.private, p.shared, p.ghost),  
    ReadWrite(p.wires)
```

```
task distribute_charge(Piece p) :  
    Reduce(p.private, p.shared, p.ghost)  
    Read(p.wires)
```

A task must declare the regions it will use & how.

*Subtask containment:
A subtask can only use (sub)regions accessible to its parent task.*

Tasks appear to execute in program order.



Execution Model

```
task simulate_circuit(Region[Node] N, Region[Wires] W) :  
{  
    ...  
    → calc_currents(piece[0], p0, s0, g0);  
    → calc_currents(piece[1], p1, s1, g1);  
    → distribute_charge(piece[0], p0, s0, g0);  
    → distribute_charge(piece[1], p1, s1, g1);  
    ...  
}
```



Interferes?
Interferes?

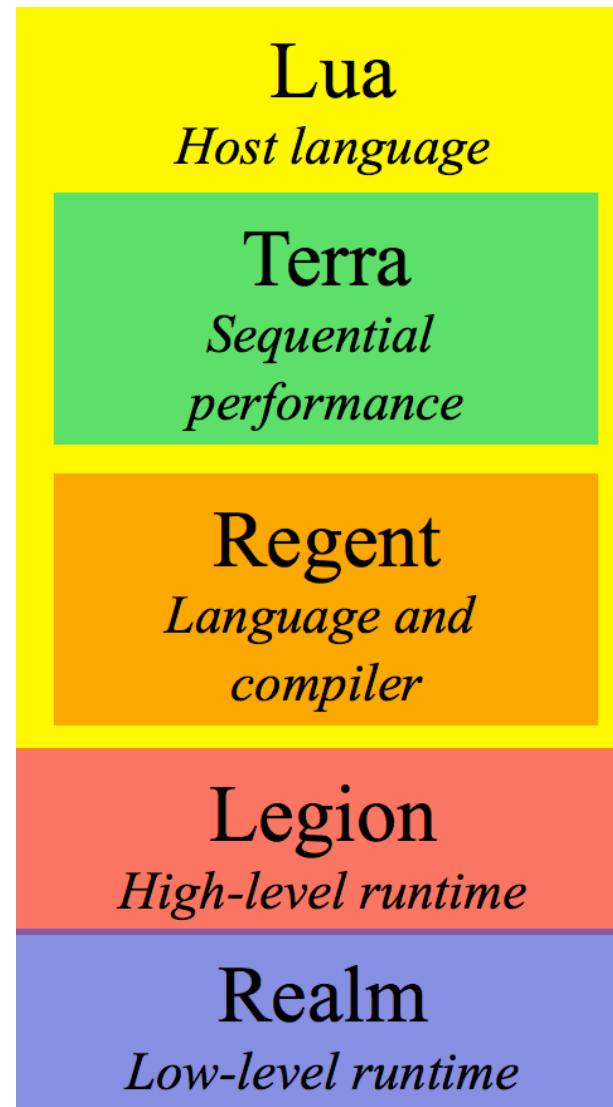
Tasks are issued in program order.



Legion Summary

- **Logical regions: a relational data model**
 - Support partitioning and slicing
 - Convey locality, independence/aliasing
- **Implicit task parallelism**
 - Task may have arbitrary sub-tasks
 - Tasks declare region usage including privileges and fields
- **Tasks appear to execute in program order**
 - Execute in parallel when non-interference established
- **Portability by separating mapping from function**

Legion as a Runtime for other Languages



Summary: Alternative Systems

PGAS: Partitioned Global Address Space (Example: UPC)

- Alternative to MPI
- Process address space split into local and global memory
- User still responsible for synchronization

Task-based models

- Only focus on actual dependencies in the code
- Critical path becomes ultimate performance metric

Two instantiations of this concept

- Charm++: actor model
- Legion: data centric / separation of concerns of specification and mapping

Such models allow for more asynchrony

- Suitable for modern heterogeneous architectures with high variability
- Challenges: grain size, scheduling, and ultimately standardization

Task-based models will likely grow in importance

Lecture Wrap-Up

Basic introduction into parallelism

- The need and purpose of using parallelism
- Architectures and application areas
- Basic patterns from master/worker and SPMD to tasking
- Metrics (Speed-Up, Efficiency)

Parallel programming APIs

- Message Passing Interface (MPI), OpenMP, Pthreads
- Alternative models: CUDA, Task-based programming, ...

Optimization and tuning

- Tuning of sequential code
- Scaling parallel code

Effective Parallelization requires rethinking of the problem

- Possibly different algorithms, data structures, ...
- Every few orders of magnitude in scale once again changes the picture

In Case You Have Further Interest in the Topic

Follow-on lectures and Praktika

- IN2106: Efficient Programming of Multicore Processors and Supercomputers
- IN2310: Parallel Program Engineering (tools, optimization)
- IN????: Parallel Programming Systems (runtime system implementations)

Watch out for courses and events at LRZ

- In particular focusing on the new SuperMUC-NG

Possibility for Master thesis, guided research, ...

- Prototyping proposals for the MPI and the OpenMP standard
- Implementation of parallel algorithms and applications
- Performance tuning and optimization
- Debugging and performance tools
- Fault tolerance
- Architecture and system evaluations
- ...

The World is Parallel – Think Parallel

