

# Getting the clowns out of the Mandelbrot

Philipp Czerner

May 7, 2018

# Overview

optimisation	time	change
baseline	1	
parallelise	0.34752	+188%
dynamic distribution	0.27227	+28%
enable optimisations	0.21961	+24%
remove complex numbers	0.07595	+189%
remove mod and branch	0.07611	+0%
loop unrolling	0.02933	+159%
main cardioid fastpath	0.00258	+1033%
fractal bulb fastpath	0.00188	+37%
distributed initialisation	0.00177	+6%
total	0.00177	×564.97

# Baseline

```
for (int i = 0; i < size_y; i++) {
    for (uint64_t j = 0; j < size_x; ++j) {
        double x = view_x0 + j * step_x;
        double y = view_y1 - (i + offset_y) * step_y;
        complex double C = x + y * I, Z = C;
        int k;
        for (k = 1; cabs(Z) < 2 && k < max_iter; ++k) {
            Z = Z * Z + C;
        }
        if (k == max_iter) {
            memcpy(image, "\0\0", 3);
        } else {
            int index = (k + palette_shift) % number_of_colors;
            memcpy(image, colors[index], 3);
        }
        image += 3;
    }
}
```

- Some tiny changes to code style and control flow, nothing performance related

## Parallelise (+188%)

```
for (int i = 0; i < work_size; ++i) {
    work_args[i] = (Mandelbrot_call_args) {(u8*)image, size_x, size_y, ... };
    work_args[i].offset_y = size_y*i / work_size;
    work_args[i].size_y = size_y*(i+1) / work_size - work_args[i].offset_y;
    work_args[i].image += work_args[i].size_x * work_args[i].offset_y * 3;
}
pthread_t threads[num_threads];
for (int i = 0; i < num_threads; ++i) {
    pthread_create(&threads[i], NULL, &mandelbrot_thread, NULL);
}
for (int i = 0; i < num_threads; ++i) {
    pthread_join(threads[i], NULL);
}
```

- Note that work is spread over threads if uneven
- Always split into contiguous chunks!
  - Here, x-axis is contiguous in memory

## Dynamic distribution (+28%)

```
#define work_size 512
Mandelbrot_call_args work_args[work_size];
atomic_int work_index;

void* mandelbrot_thread(void* _) {
    while (work_index < work_size) {
        int i = work_index++;
        if (i >= work_size) break;
        Mandelbrot_call_args* p = &work_args[i];
        mandelbrot_calculate( ... );
    }
}
```

- ▶ Synchronisation is just a single integer
- ▶ Not much to see here

## Enable optimisations (+24%)

```
__attribute__((optimize(3), __target__("avx,ssse3,sse4a")))  
void mandelbrot_calculate( ... ) { ... }
```

- Enable via an attribute as I cannot control the compilation process

## Remove complex numbers (+189%)

```
for (int j = 0; j < size_x; ++j) {
    double Cr = view_x0 + j * step_x;
    double Ci = y, Zr = 0, Zi = 0, Zr_old = 0;
    int k;
    for (k = 0; Zr*Zr + Zi*Zi < 4 && k < max_iter; ++k) {
        Zr_old = Zr;
        Zr = Zr*Zr - Zi*Zi + Cr;
        Zi = 2.0 * Zr_old * Zi + Ci;
    }
    if (k == max_iter) {
        memcpy(image, "\0\0", 3);
    } else {
        int index = (k + palette_shift) % number_of_colors;
        memcpy(image, colors[index], 3);
    }
    image += 3;
}
```

- ▶ Remove the square root inside cabs
- ▶ Write down the complex math by hand, simplify squaring a bit

## Remove mod and branch (+0%)

```
for (int j = 0; j < size_x; ++j) {  
    double Cr = view_x0 + j * step_x;  
    double Ci = y, Zr = 0, Zi = 0, Zr_old = 0;  
    int k;  
    for (k = 0; Zr*Zr + Zi*Zi < 4 && k < max_iter; ++k) {  
        Zr_old = Zr;  
        Zr = Zr*Zr - Zi*Zi + Cr;  
        Zi = 2.0 * Zr_old * Zi + Ci;  
    }  
    memcpy(image, &colors_opt[3*k], 3);  
    image += 3;  
}
```

- ▶ **Context:** Similar size of max\_iter and number\_of\_colors
- ▶ No performance gain right now, but simplifies code
- ▶ Has to be set up by the main thread beforehand (almost no overhead)



## Loop unrolling (+159%)

```
...  
for (int kk = 0; kk < max_iter; ++kk) {  
    u8 flag1 = Z1r*Z1r + Z1i*Z1i < 4;  
    u8 flag2 = Z2r*Z2r + Z2i*Z2i < 4;  
    u8 flag3 = Z3r*Z3r + Z3i*Z3i < 4;  
    u8 flag4 = Z4r*Z4r + Z4i*Z4i < 4;  
    if (!(flag1 || flag2 || flag3 || flag4)) break;  
    k1 += flag1;    k2 += flag2;    k3 += flag3;    k4 += flag4;  
    Z1r_old = Z1r; Z2r_old = Z2r; Z3r_old = Z3r; Z4r_old = Z4r;  
    Z1r = Z1r*Z1r - Z1i*Z1i + C1r;  
    Z2r = Z2r*Z2r - Z2i*Z2i + C2r;  
    Z3r = Z3r*Z3r - Z3i*Z3i + C3r;  
    Z4r = Z4r*Z4r - Z4i*Z4i + C4r;  
    Z1i = 2.0 * Z1r_old * Z1i + y;  
    Z2i = 2.0 * Z2r_old * Z2i + y;  
    Z3i = 2.0 * Z3r_old * Z3i + y;  
    Z4i = 2.0 * Z4r_old * Z4i + y;  
}  
...
```

- Manually unroll the loop to process four elements at once.

## Loop unrolling (+159%)

- ▶ No branches for the finished elements!
  - ▶ not trivial to make work, compiler cannot do this
  - ▶ if  $|Z| > 2$  in some iteration, it will continue to be
- ▶ `size_x` may not be a multiple of 4
  - ▶ must process remaining elements afterwards

# Manual vectorisation

```
typedef double   v4d __attribute__((vector_size (sizeof(double)  *4)));
typedef long int v4i __attribute__((vector_size (sizeof(long int)*4)));
v4d Cr;
for (int m = 0; m < 4; ++m)
    Cr[m] = view_x0 + (j+m) * step_x;
double Ci = y;
v4d Zr = {0}, Zi = {0}, Zr_old = {0};
v4i k = {0};
for (int kk = 0; kk < max_iter; ++kk) {
    v4i tmp = (Zr*Zr + Zi*Zi) < 4;
    k -= tmp; // Note that vector comparison set all bits
    if (!(tmp[0] || tmp[1] || tmp[2] || tmp[3])) break;
    Zr_old = Zr;
    Zr = Zr*Zr - Zi*Zi + Cr;
    Zi = 2.0 * Zr_old * Zi + y;
}
for (int m = 0; m < 4; ++m)
    memcpy(image+3*m, &colors_opt[3*k[m]], 3);
```

► Nice and easy to use

# Manual vectorisation

```
typedef double   v4d __attribute__((vector_size (sizeof(double)  *4)));
typedef long int v4i __attribute__((vector_size (sizeof(long int)*4)));
v4d Cr;
for (int m = 0; m < 4; ++m)
    Cr[m] = view_x0 + (j+m) * step_x;
double Ci = y;
v4d Zr = {0}, Zi = {0}, Zr_old = {0};
v4i k = {0};
for (int kk = 0; kk < max_iter; ++kk) {
    v4i tmp = (Zr*Zr + Zi*Zi) < 4;
    k -= tmp; // Note that vector comparison set all bits
    if (!(tmp[0] || tmp[1] || tmp[2] || tmp[3])) break;
    Zr_old = Zr;
    Zr = Zr*Zr - Zi*Zi + Cr;
    Zi = 2.0 * Zr_old * Zi + y;
}
for (int m = 0; m < 4; ++m)
    memcpy(image+3*m, &colors_opt[3*k[m]], 3);
```

- ▶ Nice and easy to use
- ▶ Makes things about 12% slower!


## Domain-specific optimisations

- ▶ Done with routine tricks, now the real work starts
- ▶ Bottleneck is in inner loop, CPU-bound
  - ▶ Use profiler / custom timing code (see below) to verify
- ▶ Try to figure out behaviour of the program, have to look at what the code is actually trying to accomplish

```
// Queries high-precision timer. Not necessarily related to actual  
// wall-clock time, use OS facilities for that.
```


```
uint64_t rdtsc() {  
    uint32_t l, h;  
    asm volatile ("rdtscp\nmov %%edx, %0\nmov %%eax, %1\n\t":  
                  "=r" (h), "=r" (l) :: "%eax", "%edx");  
    return (uint64_t)h << 32 | (uint64_t)l;  
}
```

# Domain-specific optimisations



- ▶ Where is the time spent?
  - ▶ No need to guess, just measure.

## Domain-specific optimisations



- ▶ White is more time
- ▶ The main cardioid is about a third, and quite expensive


## Main cardioid fastpath (+1033%)

```
u64 j;
for (j = 0; j+4 <= size_x; j += 4) {
    while (j < size_x) {
        double a = view_x0 + j*step_x - 0.25;
        double q = a*a + y*y;
        if (q*(q+a) >= y*y * 0.25) break;
        memcpy(image, "\0\0", 3);
        ++j;
        image += 3;
    }
    if (j+4 > size_x) break;
    ...
}
```

- ▶ **Context:** This optimisation is only good when the main cardioid makes up a large part of the image!
- ▶ Use formula to detect the inside (I simply looked it up)
- ▶ Check is quite cheap, can do it on every pixel
- ▶ Same thing for the bulb on the top (additional +37%)



Main cardioid fastpath (+1033%)



► Much better.

## Distributed initialisation (+6%)

```
int thread_count;
pthread_t threads[64];
int work_size;
Mandelbrot_call_args work_args[1000];
atomic_int thread_index, work_index;

void* mandelbrot_thread(void* _) {
    while (thread_index < thread_count) {
        int i = thread_index++;
        if (i >= thread_count) break;
        pthread_create(&threads[i], NULL, &mandelbrot_thread, NULL);
    }
    while (work_index < work_size) {
        int i = work_index++;
        if (i >= work_size) break;
        Mandelbrot_call_args* p = &work_args[i];
        mandelbrot_calculate( ... );
    }
}
```

- Let the threads initialise each other

# Final notes

- ▶ Don't blindly trust the compiler to do non-trivial optimisations—you know much better which transformations are safe
- ▶ Look at the actual data, effort spent visualising is seldom wasted
- ▶ Code is now fast enough ( $\sim 10\text{ms}$ ) that constant overheads (initialisation, synchronisation) become relevant again
- ▶ No changes to the calculation (e.g. low precision approximation or such) as Mandelbrot is sensitive to rounding errors