

Pruning the clown sheep of the family





Philipp Czerner

Remark on the previous assignments

```
pthread_attr_t attrs;
pthread_attr_init(&attrs);
int main_cpu = sched_getcpu();
int num_cores = sysconf(_SC_NPROCESSORS_ONLN);
for (int i = 0; i < num_threads; ++i) {
    int thread_cpu = (main_cpu + i) % num_cores;
    cpu_set_t cpu;
    CPU_ZERO(&cpu);
    CPU_SET(thread_cpu, &cpu);
    pthread_attr_setaffinity_np(&attrs, sizeof(cpu_set_t), &cpu);
    pthread_create(&threads[i], &attrs, thread_function, &args[i]);
}
pthread_attr_destroy(&attrs);
```

- ▶ Last time, OpenMP was still faster than even my crazy initialisation
- ▶ A large part of that is setting CPU affinity

Overview

optimisation	time	change	difficulty	generality
baseline	1			
precomputed lookup table	0.00000354	281908		
OpenMP	0.00028949	0.0123		
total	0.00000354	281908		

- Baseline already has optimisations enabled, although only for my functions

Baseline

```
void traverse_rec(tree* node) {
    if (node != NULL) {
        node->IQ = compute_IQ(node->data);
        genius[node->id] = node->IQ;
        traverse_rec(node->left );
        traverse_rec(node->right);
    }
}

int is_prime_or_zero_or_one(int n) {
    for (int i = 2; i <= n/2; ++i)
        if (n % i == 0) return false;
    return true;
}

int compute_IQ(int data) {
    int sum = 0;
    for (int i = 0; i < data; i++)
        sum += is_prime_or_zero_or_one(i);
    return 70 + (sum % 100)*(sum % 100)*(sum % 100)*(sum % 100)/1000000;
}
```

- The code of compute_IQ will be irrelevant

Precomputed lookup table ($\times 281908$)

```
void traverse_rec(tree* node) {  
    if (node != NULL) {  
        node->IQ = lookup_iq[node->data];  
        genius[node->id] = node->IQ;  
        traverse_rec(node->left );  
        traverse_rec(node->right);  
    }  
}
```

- **Context:** There are only 9113 possible values for node->data
- **Context:** The function compute_IQ easily takes more than 1000 cycles to compute, even if optimised

Working inside arbitrary restrictions

- ▶ Computing the lookup at runtime is unnecessary, takes too long
- ▶ Instead, store it in the binary
- ▶ Problem?

Working inside arbitrary restrictions

- ▶ Computing the lookup at runtime is unnecessary, takes too long
- ▶ Instead, store it in the binary
- ▶ Problem? We have to fit it into the 10 KiB limit
- ▶ Fitting 9113 values into 10240 bytes \Rightarrow Compress data
- ▶ How? Just look at it.

Working inside arbitrary restrictions

[illegible]

The glorious C preprocessor

- ▶ Lots of numbers repeat!
- ▶ Length of repetitions only increases later on
- ▶ Compact encoding using the C preprocessor
- ▶ `R4(7)` expands to `7,7,7,7`

```
#define R2(x) x,x
#define R3(x) x,x,x
#define R4(x) R3(x),x
#define R6(x) R4(x),x,x
#define R8(x) R6(x),x,x
...
#define R44(x) R42(x),x,x
#define R52(x) R44(x),R8(x)
#define R104(x) R52(x),R52(x)
#define R208(x) R104(x),R104(x)
```

The glorious C preprocessor

```
uint8_t iq_lookup[9113] = {R104(70),R10(70),R38(71),R22(72),R18(73),R8(74),  
    R24(75),R6(76),R10(77),R2(78),R16(79),R6(80),R6(81),R8(82),R4(83),R2(84),  
    R10(85),R14(86),R4(87),R2(88),R4(90),R14(91),R6(92),R10(94),R2(95),R4(96),  
    R6(98),R8(99),R6(101),R6(103),R4(105),R6(107),R8(108),R4(110),R8(113),  
    R10(115),R2(117),R10(119),R2(122),R6(124),R4(127),R6(129),R8(132), ... };
```

- ▶ ~4800 bytes of source code
- ▶ Generated by a small Python script
 - ▶ Also optimises the `#defines` at the beginning

OpenMP ($\times 0.0123$)

```
void traverse_rec(tree* node) {  
    if (node != NULL) {  
        node->IQ = iq_lookup[node->data];  
        genius[node->id] = node->IQ;  
  
        #pragma omp task  
        traverse_rec(node->left );  
  
        traverse_rec(node->right);  
    }  
}
```

- ▶ Could be optimised further, but has no chance of beating the previous version
 - ▶ Creating a thread alone takes quadruple of its total time

Questions?

For later, `noclowns@nicze.de` or just talk to me.